

Lecture 1: Intro to logistic regression

Agenda

- Introductions
- Overview of course details
- Begin logistic regression
- HW1 released on course website

Class overview

- STA 711 focuses on *statistical inference*: estimation, confidence intervals, and hypothesis testing
- Throughout the semester, topics will be initially motivated by logistic regression
- We will continue with inference and GLMs in STA 712 (Generalized Linear Models)

Grading philosophy

- Focusing on grades can detract from the learning process
- Homework should be an opportunity to *practice* the material. It is ok to make mistakes when practicing, as long as you make an honest effort
- Errors are a good opportunity to learn and revise your work
- Partial credit and weighted averages of scores make the meaning of a grade confusing. Does an 85 in the course mean you know 85% of everything, or everything about 85% of the material?

Grading in this course

- I will give you feedback on every assignment
- All assignments are graded as Mastered / Not yet mastered
- If you haven't yet mastered something, you get to try again!

Course components

- Regular homework assignments
 - Practice material from class
 - You may resubmit “Not yet mastered” questions once
- 3 take-home exams
 - Opportunity to demonstrate mastery of course material
 - Optional make-up exams for “Not yet mastered” questions
- Optional final exam
 - Final opportunity to demonstrate mastery

Assigning grades

To get an **A-** in the course:

- Receive credit for at least $N-2$ homework assignments
- Master at least 80% of the questions on all three exams

To get an **A** in the course:


- Receive credit for at least $N-1$ homework assignments
- Master at least 80% of the questions on all three exams

Late work and resubmissions

- You get a bank of **5** extension days. You can use 1–2 days on any assignment, exam, or project.
- No other late work will be accepted (except in extenuating circumstances!)

Motivating example: Dengue fever

Dengue fever: a mosquito-borne viral disease affecting 400 million people a year



Motivating example: Dengue data

Data: Data on 5720 Vietnamese children, admitted to the hospital with possible dengue fever. Variables include:

- *Sex*: patient's sex (female or male)
- *Age*: patient's age (in years)
- *WBC*: white blood cell count
- *PLT*: platelet count
- other diagnostic variables...
- *Dengue*: whether the patient has dengue (0 = no, 1 = yes)

Motivating example: Dengue data

Research questions:

- How well can we predict whether a patient has dengue?
- Which diagnostic measurements are most useful?
- Is there a significant relationship between WBC and dengue?

Research questions

- How well can we predict whether a patient has dengue?
- Which diagnostic measurements are most useful?
- Is there a significant relationship between WBC and dengue?

How can I answer each of these questions? Discuss with a neighbor for 2 minutes, then we will discuss as a class.

Fitting a model: initial attempt

What if we try a linear regression model?

Y_i = dengue status of i th patient

$$Y_i = \beta_0 + \beta_1 \text{WBC}_i + \varepsilon_i \quad \varepsilon_i \stackrel{\text{iid}}{\sim} N(0, \sigma_\varepsilon^2)$$

What are some potential issues with this linear regression model?

Second attempt

Let's rewrite the linear regression model:


Second attempt

$$Y_i \sim \text{Bernoulli}(p_i) \quad p_i = \mathbb{P}(Y_i = 1 | \text{WBC}_i)$$

$$p_i = \beta_0 + \beta_1 \text{WBC}_i$$

Are there still any potential issues with this approach?

Don't fit linear regression with a binary response



Fixing the issue: logistic regression

$$Y_i \sim \text{Bernoulli}(p_i)$$

$$g(p_i) = \beta_0 + \beta_1 \text{WBC}_i$$

where $g : (0, 1) \rightarrow \mathbb{R}$ is unbounded.

Usual choice: $g(p_i) = \log\left(\frac{p_i}{1 - p_i}\right)$

Odds

Definition: If $p_i = \mathbb{P}(Y_i = 1 | WBC_i)$, the **odds** are $\frac{p_i}{1 - p_i}$

Example: Suppose that $\mathbb{P}(Y_i = 1 | WBC_i) = 0.8$. What are the *odds* that the patient has dengue?

Odds

Definition: If $p_i = \mathbb{P}(Y_i = 1 | WBC_i)$, the **odds** are $\frac{p_i}{1 - p_i}$

The probabilities $p_i \in [0, 1]$. The linear function $\beta_0 + \beta_1 WBC_i \in (-\infty, \infty)$. What range of values can $\frac{p_i}{1 - p_i}$ take?

Log odds

$$g(p_i) = \log\left(\frac{p_i}{1 - p_i}\right)$$

Binary logistic regression

$$Y_i \sim \text{Bernoulli}(p_i)$$

$$\log\left(\frac{p_i}{1 - p_i}\right) = \beta_0 + \beta_1 \text{WBC}_i$$

Note: Can generalize to $Y_i \sim \text{Binomial}(m_i, p_i)$, but we won't do that yet.

Example: simple logistic regression

Y_i = dengue status (0 = no, 1 = yes) $Y_i \sim \text{Bernoulli}(p_i)$

$$\log\left(\frac{\hat{p}_i}{1 - \hat{p}_i}\right) = 1.737 - 0.361 \text{ WBC}_i$$

Work in groups of 2-3 for 5 minutes on the following questions:

- Are patients with a higher WBC more or less likely to have dengue?
- Interpret the estimated slope in context of a unit change in the log odds.
- What is the change in *odds* associated with a unit increase in WBC?

